

Buildings built in minutes - An SfM Approach

This article is written by Chahat Deep Singh. If you have any questions/corrections regarding the article, please email at [chahat\[at\]terpmail\[dot\]umd\[edu\]](mailto:chahat[at]terpmail[dot]umd[edu]).

To be submitted in a group of two.

Table of Contents:

- 1. Deadline
- 2. Introduction
- 3. Phase 1: Traditional Approach
 - 3.1. Feature Matching
 - 3.2. Estimating Fundamental Matrix
 - 3.2.1. Epipolar Geometry
 - 3.2.2. Fundamental Matrix
 - 3.2.3. Match Outlier Rejection using RANSAC
 - 3.3. Estimate *Essential Matrix* from Fundamental Matrix
 - 3.4. Estimate Camera Pose from Essential Matrix
 - 3.5. Check for Cheirality Condition using Triangulation
 - 3.5.1. Non-Linear Triangulation
 - 3.6. Perspective- n -points
 - 3.6.1. Linear Camera Pose Estimation
 - 3.6.2. PnP RANSAC
 - 3.6.3. NonLinear PnP
 - 3.7. Bundle Adjustment
 - 3.7.1. Visibility Matrix
 - 3.7.2. Bundle Adjustment
- 4. Putting the pipeline together
- 5. Notes about Data Set
- 6. Submission Guidelines
 - 6.1. File tree and naming
 - 6.2. Report
- 7. Collaboration Policy

1. Deadline

11:59PM, April 28, 2024.

2. Introduction

We have been playing with images for so long, mostly in 2D scene. Recall [project 1](#) where we stitched multiple images with about 30-50% common features between a couple of images. Now let's learn how to **reconstruct a 3D scene and simultaneously obtain the camera poses** of a monocular camera w.r.t. the given scene. This procedure is known as Structure from Motion (SfM). As the name suggests, you are creating the entire **rigid** structure from a set of images with different view points (or equivalently a camera in motion). A few years ago, Agarwal et. al published [Building Rome in a Day](#) in which they reconstructed the entire city just by using a large collection of photos from the Internet. Ever heard of Microsoft [Photosynth](#)? *Facinating? isn't it!?* There are a few open source SfM algorithm available online like [VisualSFM](#). *Try them!*

Let's learn how to recreate such algorithm. There are a few steps that collectively form SfM:

- **Feature Matching** and Outlier rejection using **RANSAC**
- Estimating **Fundamental Matrix**
- Estimating **Essential Matrix** from Fundamental Matrix
- Estimate **Camera Pose** from Essential Matrix
- Check for **Cheirality Condition** using **Triangulation**
- **Perspective-n-Point**
- **Bundle Adjustment**

3. Traditional Approach to the SfM problem

3.1. Feature Matching, Fundamental Matrix and RANSAC

We have already learned about keypoint matching using SIFT keypoints and descriptors (Recall Project 1: Panorama Stitching). It is important to refine the matches by rejecting outlier correspondence. Before rejecting the correspondences, let us first understand what Fundamental matrix is!




Figure 1: Feature matching between two images from different views.

3.2. Estimating Fundamental Matrix

The fundamental matrix, denoted by F , is a 3×3 (*rank 2*) matrix that relates the corresponding set of points in two images from different views (or stereo images). But in order to understand what fundamental matrix actually is, we need to understand what *epipolar geometry* is! The epipolar geometry is the intrinsic projective geometry between two views. It only depends on the cameras' internal parameters (K matrix) and the relative pose *i.e.* it is **independent of the scene structure**.

3.2.1. Epipolar Geometry

Let's say a point \mathbf{X} in the 3D-space (viewed in two images) is captured as \mathbf{x} in the first image and \mathbf{x}' in the second. *Can you think how to formulate the relation between the corresponding image points \mathbf{x} and \mathbf{x}' ?* Consider Fig. 2. Let \mathbf{C} and \mathbf{C}' be the respective camera centers which forms the baseline for the stereo system. Clearly, the points \mathbf{x} , \mathbf{x}' and \mathbf{X} (or \mathbf{C} , \mathbf{C}' and \mathbf{X}) are coplanar *i.e.* $\overrightarrow{\mathbf{C}\mathbf{x}} \cdot \left(\overrightarrow{\mathbf{C}\mathbf{C}'} \times \overrightarrow{\mathbf{C}'\mathbf{x}'} \right) = 0$ and the plane formed can be denoted by π . Since these points are coplanar, the rays back-projected from \mathbf{x} and \mathbf{x}' intersect at \mathbf{X} . This is the most significant property in searching for a correspondence.




Figure 2(a): Caption goes here.




Figure 2(b): Caption goes here.

Now, let us say that only \mathbf{x} is known, not \mathbf{x}' . We know that the point \mathbf{x}' lies in the plane π which is governed by the camera baseline $\mathbf{CC'}$ and $\overrightarrow{\mathbf{Cx}}$. Hence the point \mathbf{x}' lies on the line of intersection of $\mathbf{l'}$ of π with the second image plane. The line $\mathbf{l'}$ is the image in the second view of the ray back-projected from \mathbf{x} . This line $\mathbf{l'}$ is called the *epipolar line* corresponding to \mathbf{x} . The benefit is that you don't need to search for the point corresponding to \mathbf{x} in the entire image plane as it can be restricted to the $\mathbf{l'}$.

- **Epipole** is the point of intersection of the line joining the camera centers with the image plane. (see \mathbf{e} and $\mathbf{e'}$ in the Fig. 2(a))
- **Epipolar plane** is the plane containing the baseline.
- **Epipolar line** is the intersection of an epipolar plane with the image plane. *All the epipolar lines intersect at the epipole.*

3.2.2. The Fundamental Matrix \mathbf{F}

The \mathbf{F} matrix is only an algebraic representation of epipolar geometry and can both geometrically (*constructing the epipolar line*) and arithmetically. (See [derivation \(Page 242\)](#)) ([Fundamental Matrix Song](#)) As a result, we obtain: $\mathbf{x}_i'^T \mathbf{F} \mathbf{x}_i = 0$ where $i = 1, 2, \dots, m$. This is known as epipolar constraint or correspondance condition (or *Longuet-Higgins* equation). Since, \mathbf{F} is a 3×3 matrix, we can set up a homogenous linear system with 9 unknowns:

$$\begin{bmatrix} x_i' & y_i' & 1 \end{bmatrix} \begin{bmatrix} f_{11} & f_{12} & f_{13} \\ f_{21} & f_{22} & f_{23} \\ f_{31} & f_{32} & f_{33} \end{bmatrix} \begin{bmatrix} x_i \\ y_i \\ 1 \end{bmatrix} = 0$$

$$x_i x_i' f_{11} + x_i y_i' f_{21} + x_i f_{31} + y_i x_i' f_{12} + y_i y_i' f_{22} + y_i f_{32} + x_i' f_{13} + y_i' f_{23} + f_{33} = 0$$

Simplifying for m correspondences,

$$\begin{bmatrix} x_1 x_1' & x_1 y_1' & x_1 & y_1 x_1' & y_1 y_1' & y_1 & x_1' & y_1' & 1 \\ \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots \\ x_m x_m' & x_m y_m' & x_m & y_m x_m' & y_m y_m' & y_m & x_m' & y_m' & 1 \end{bmatrix} \begin{bmatrix} f_{11} \\ f_{21} \\ f_{31} \\ f_{12} \\ f_{22} \\ f_{32} \\ f_{13} \\ f_{23} \\ f_{33} \end{bmatrix} = 0$$

How many points do we need to solve the above equation? Think! Twice! Remember *homography*, where each point correspondence contributes two constraints? Unlike homography, in \mathbf{F} matrix estimation, each point only contributes one constraints as the epipolar constraint is a scalar equation. Thus, we require at least 8 points to solve the above homogenous system. That is why it is known as [Eight-point algorithm](#).

With $N \geq 8$ correspondences between two images, the fundamental matrix, F can be obtained as: By stacking the above equation in a matrix A , the equation $Ax = 0$ is obtained. This system of equation can be answered by solving the linear least squares using Singular Value Decomposition (SVD) as explained in the [Math module](#). When applying SVD to matrix \mathbf{A} , the decomposition \mathbf{USV}^T would be obtained with \mathbf{U} and \mathbf{V} orthonormal matrices and a diagonal matrix \mathbf{S} that contains the singular values. The singular values σ_i where $i \in [1, 9], i \in \mathbb{Z}$, are positive and are in decreasing order with $\sigma_9 = 0$ since we have 8 equations for 9 unknowns.

Thus, the last column of \mathbf{V} is the true solution given that $\sigma_i \neq 0 \forall i \in [1, 8], i \in \mathbb{Z}$. However, due to noise in the correspondences, the estimated \mathbf{F} matrix can be of rank 3 *i.e.* $\sigma_9 \neq 0$. So, to enforce the rank 2 constraint, the last singular value of the estimated \mathbf{F} must be set to zero. If F has a full rank then it will have an empty null-space *i.e.* it won't have any point that is on entire set of lines. Thus, there wouldn't be any epipoles. See Fig. 3 for full rank comparisons for F matrices.




Figure 3: F Matrix: Rank 3 vs Rank 2 comparison

In MATLAB, you can use `svd` to solve $\mathbf{Ax} = 0$

```
[U, S, V] = svd(A);  
x = V(:, end);  
F = reshape(x, [3,3])';
```

To summarize, write a function `EstimateFundamentalMatrix.py` that linearly estimates a fundamental matrix F , such that $x_2^T F x_1 = 0$. The fundamental matrix can be estimated by solving the linear least squares $Ax = 0$.

3.2.3. Match Outlier Rejection via RANSAC

Since the point correspondences are computed using SIFT or some other feature descriptors, the data is bound to be noisy and (in general) contains several outliers. Thus, to remove these outliers, we use RANSAC algorithm (*Yes! The same as used in Panorama stitching!*) to obtain a

better estimate of the fundamental matrix. So, out of all possibilities, the \mathbf{F} matrix with maximum number of inliers is chosen. Below is the pseduo-code that returns the \mathbf{F} matrix for a set of matching corresponding points (computed using SIFT) which maximizes the number of inliers.

```

n=0;
for i = 1:M do
    // Choose 8 correspondences,  $\hat{x}_1$  and  $\hat{x}_2$  randomly
    F = EstimateFundmentalMatrix( $\hat{x}_1$ ,  $\hat{x}_2$ );
    S =  $\emptyset$ ;
    for j = 1:N do
        if  $|x_{2j}^T F x_{1j}| < \epsilon$  then
            S = S  $\cup$  {j}
        end
    end
    if  $n < |S|$  then
        n = |S|;
        Sin = S
    end
end
end

```

Algorithm 1: Get Inliers RANSAC




Figure 4: Feature matching after RANSAC. (Green: Selected correspondences; Red: Rejected correspondences)

Given, $N \geq 8$ correspondences between two images, $x_1 \leftrightarrow x_2$, implement a function `GetInlierRANSANC.py` that estimates inlier correspondences using fundamental matrix based RANSAC.

3.3. Estimate *Essential Matrix* from Fundamental Matrix

Since we have computed the \mathbf{F} using epipolar constraints, we can find the relative camera poses between the two images. This can be computed using the *Essential Matrix*, \mathbf{E} . Essential matrix is another 3×3 matrix, but with some additional properties that relates the corresponding points assuming that the cameras obey the pinhole model (unlike \mathbf{F}). More specifically, $\mathbf{E} = \mathbf{K}^T \mathbf{F} \mathbf{K}$ where \mathbf{K} is the camera calibration matrix or camera intrinsic matrix. Clearly, the essential matrix can be extracted from \mathbf{F} and \mathbf{K} . As in the case of \mathbf{F} matrix computation, the singular values of \mathbf{E} are not necessarily $(1, 1, 0)$ due to the noise in \mathbf{K} . This can be corrected

by reconstructing it with $(1, 1, 0)$ singular values, i.e. $\mathbf{E} = U \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix} V^T$

It is important to note that the \mathbf{F} is defined in the original image space (i.e. pixel coordinates) whereas \mathbf{E} is in the normalized image coordinates. Normalized image coordinates have the origin at the optical center of the image. Also, relative camera poses between two views can be computed using \mathbf{E} matrix. Moreover, \mathbf{F} has 7 degrees of freedom while \mathbf{E} has 5 as it takes camera parameters in account. ([5-Point Motion Estimation Made Easy](#))

Given F , estimate the essential matrix $E = K^T F K$ by implementing the function `EssentialMatrixFromFundamentalMatrix.py`.

3.4. Estimate **Camera Pose** from Essential Matrix

The camera pose consists of 6 degrees-of-freedom (DOF) Rotation (Roll, Pitch, Yaw) and Translation (X, Y, Z) of the camera with respect to the world. Since the \mathbf{E} matrix is identified, the four camera pose configurations: $(C_1, R_1), (C_2, R_2), (C_3, R_3)$ and (C_4, R_4) where $C \in \mathbb{R}^3$ is the camera center and $R \in SO(3)$ is the rotation matrix, can be computed. Thus, the camera pose can be written as: $P = KR [I_{3 \times 3} \quad -C]$ These four pose configurations

can be computed from \mathbf{E} matrix. Let $\mathbf{E} = UDV^T$ and $W = \begin{bmatrix} 0 & -1 & 0 \\ 1 & 0 & 0 \\ 0 & 0 & 1 \end{bmatrix}$. The four configurations can be written as:

1. $C_1 = U(:, 3)$ and $R_1 = UWV^T$
2. $C_2 = -U(:, 3)$ and $R_2 = UWV^T$
3. $C_3 = U(:, 3)$ and $R_3 = UW^T V^T$
4. $C_4 = -U(:, 3)$ and $R_4 = UW^T V^T$

It is important to note that the $\det(R) = 1$. If $\det(R) = -1$, the camera pose must be corrected *i.e.* $C = -C$ and $R = -R$.

Implement the function `ExtractCameraPose.py`, given E .

3.5. Triangulation Check for Cheirality Condition

In the previous section, we computed four different possible camera poses for a pair of images using essential matrix. In this section we will triangulate the 3D points, given two camera poses.

Given two camera poses, (C_1, R_1) and (C_2, R_2) , and correspondences, $x_1 \leftrightarrow x_2$, triangulate 3D points using linear least squares. Implement the function `LinearTriangulation.py`.

Though, in order to find the *correct* unique camera pose, we need to remove the disambiguity. This can be accomplished by checking the **cheirality condition** *i.e. the reconstructed points must be in front of the cameras*. To check the cheirality condition, triangulate the 3D points (given two camera poses) using **linear least squares** to check the sign of the depth Z in the camera coordinate system w.r.t. camera center. A 3D point X is in front of the camera iff: $r_3(X - C) > 0$ where r_3 is the third row of the rotation matrix (z-axis of the camera). Not all triangulated points satisfy this condition due to the presence of correspondence noise. The best camera configuration, (C, R, X) is the one that produces the maximum number of points satisfying the cheirality condition.




Figure 5: Initial triangulation plot with disambiguation, showing all four possible camera poses.

Given four camera pose configurations and their triangulated points, find the unique camera pose by checking the cheirality condition - the reconstructed points must be in front of the cameras (implement the function `DisambiguateCameraPose.py`).

3.5.1. Non-Linear Triangulation

Given two camera poses and linearly triangulated points, \tilde{X} , the locations of the 3D points that minimizes the reprojection error (Recall [Project 2](#)) can be refined. The linear triangulation minimizes the algebraic error. Though, the reprojection error is geometrically meaningful error and can be computed by measuring error between measurement and projected 3D point:

$$\min_x \sum_{j=1,2} \left(u^j - \frac{P_1^{jT} \tilde{X}}{P_3^{jT} \tilde{X}} \right)^2 + \left(v^j - \frac{P_2^{jT} \tilde{X}}{P_3^{jT} \tilde{X}} \right)^2$$

Here, j is the index of each camera, \widetilde{X} is the homogeneous representation of X . P_i^T is each row of camera projection matrix, P . This minimization is highly nonlinear due to the divisions. The initial guess of the solution, X_0 , is estimated via the linear triangulation to minimize the cost function. This minimization can be solved using nonlinear optimization functions such as `scipy.optimize.leastsq` or `scipy.optimize.least_squares` in Scipy library.

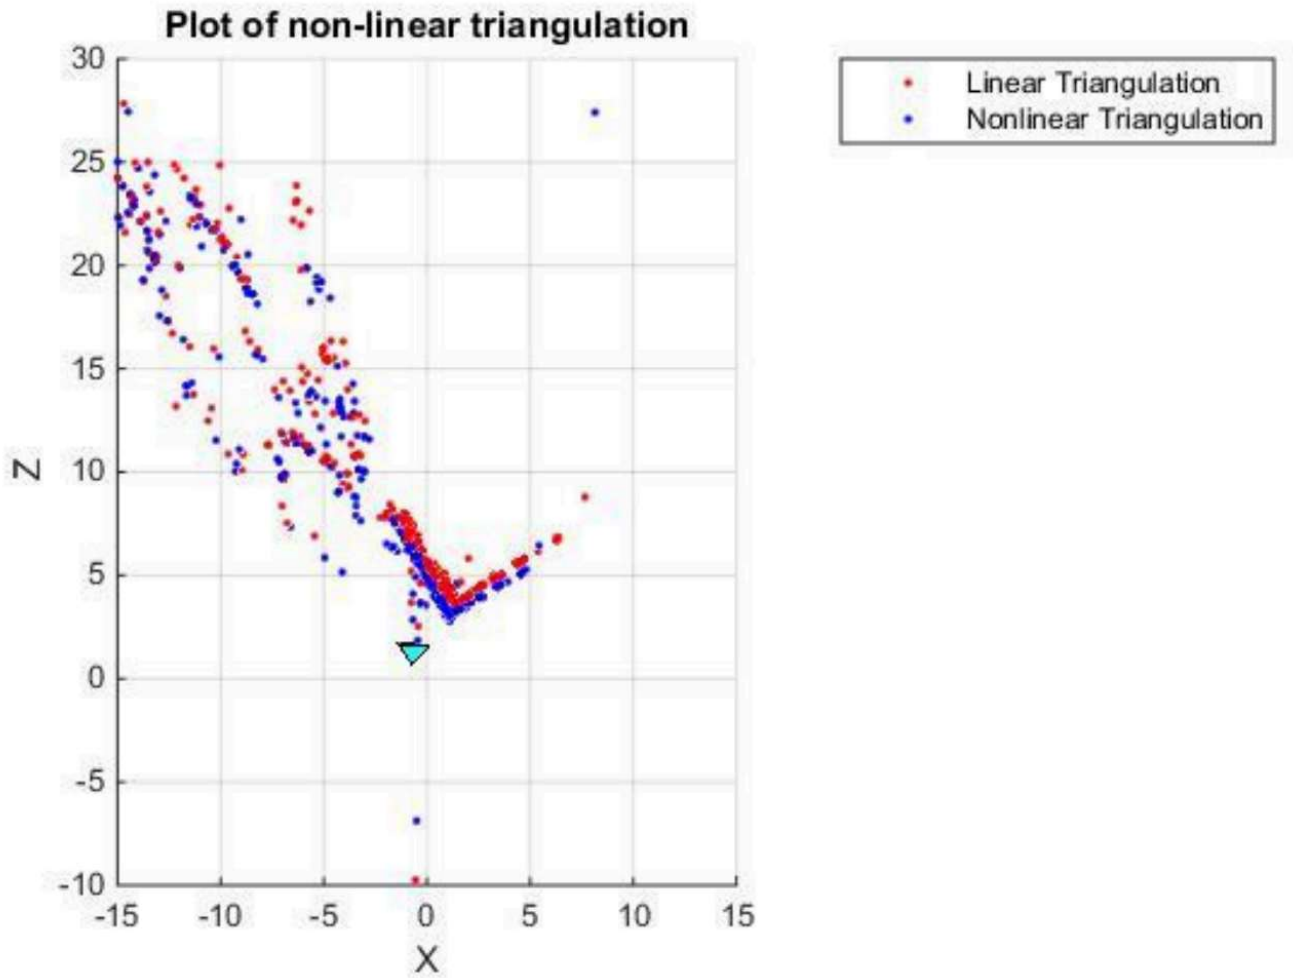


Figure 6: Comparison between non-linear vs linear triangulation.

Given two camera poses and linearly triangulated points, X , refine the locations of the 3D points that minimizes reprojection error (implement the function `NonlinearTriangulation.py`).

3.6. Perspective- n -Points

Now, since we have a set of n 3D points in the world, their $2D$ projections in the image and the intrinsic parameter; the 6 DOF camera pose can be estimated using linear least squares. This fundamental problem, in general is known as *Persepective- n - Point* (PnP). For there to exist a

solution, $n \geq 3$. There are multiple methods to solve the PnP problem and have an assumptions in most of them that the camera is calibrated. Methods such as [Unified PnP](#) (or UPnP) do not abide with the said assumption as they estimate both intrinsic and extrinsic parameters. In this section, you will a simpler version of PnP. You will register a new image given 2D-3D correspondences, i.e. $X \leftrightarrow x$ followed by nonlinear optimization.

3.6.1 Linear Camera Pose Estimation

Given 2D-3D correspondences, $X \leftrightarrow x$ and the intrinsic parameter K , estimate the camera pose using linear least squares (implement the function `LinearPnP.py`). 2D points can be normalized by the intrinsic parameter to isolate camera parameters, (C, R) , i.e. $K^{-1}x$. A linear least squares system that relates the 3D and 2D points can be solved for (t, R) where $t = -R^T C$. Since the linear least square solve does not enforce orthogonality of the rotation matrix, $R \in SO(3)$, the rotation matrix must be corrected by $R = UV^T$ where $R = UDV^T$. If the corrected rotation has -1 determinant, $R = -R$. This linear PnP requires at least 6 correspondences. (*Think why?*)

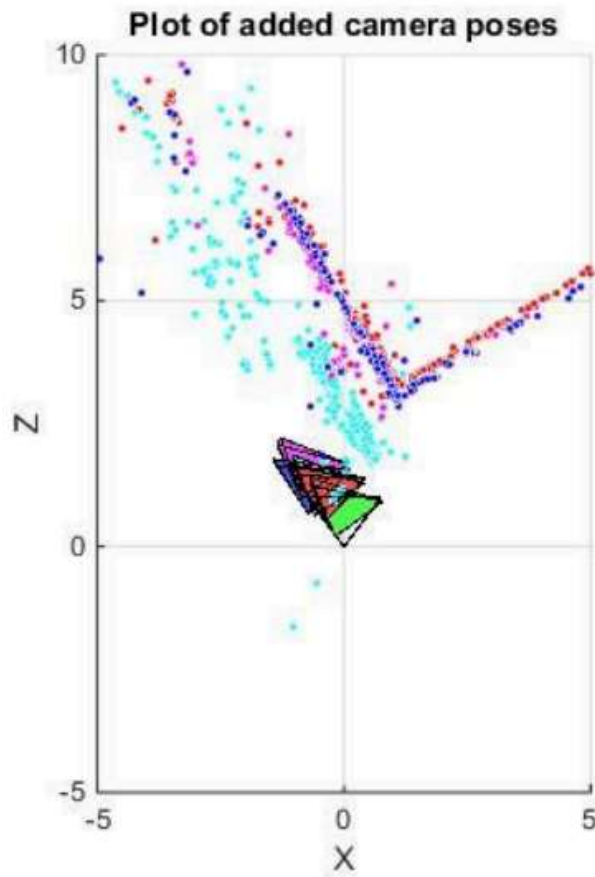


Figure: Plot of the camera poses with feature points. Different color represents feature correspondences from different pair of images. Blue points are features from Image 1 and Image 2; Red points are features from Image 2 and Image 3 etc.

3.6.2 PnP RANSAC

PnP is prone to error as there are outliers in the given set of point correspondences. To overcome this error, we can use RANSAC (yes, again!) to make our camera pose more robust to outliers. To formalize, given $N \geq 6$ 3D-2D correspondences, $X \leftrightarrow x$, implement the following function that estimates camera pose (C, R) via RANSAC (implement the function PnP_RANSAC.py).

The algorithm below depicts the solution with RANSAC.

```
n = 0
for i = 1:M do
    // Choose 6 correspondences,  $\hat{X}$  and  $\hat{x}$ , randomly
    [C R] = LinearPnP( $\hat{X}$ ,  $\hat{x}$ , K);
    S =  $\emptyset$ ;
    for j = 1:N do
        // Measure Reprojection error
        
$$e = \left(u - \frac{P_1^T \tilde{X}}{P_3^T \tilde{X}}\right)^2 + \left(v - \frac{P_2^T \tilde{X}}{P_3^T \tilde{X}}\right)^2;$$

        if  $e < \epsilon_r$  then
            | S = S  $\cup$  {j}
        end
    end
    if  $n < |S|$  then
        |  $n = |S|$ ;
        |  $S_{in} = S$ 
    end
end
```

Algorithm 2: PnP RANSAC

Just like in triangulation, since we have the linearly estimated camera pose, we can refine the camera pose that minimizes the reprojection error (Linear PnP only minimizes the algebraic error).

3.6.3 Nonlinear PnP

Given $N \geq 6$ 3D-2D correspondences, $X \leftrightarrow x$, and linearly estimated camera pose, (C, R) , refine the camera pose that minimizes reprojection error (implement the function `NonlinearPnP.py`). The linear PnP minimizes algebraic error. Reprojection error that is geometrically meaningful error is computed by measuring error between measurement and projected 3D point

$$\min_{C, R} \sum_{i=1, J} \left(u^j - \frac{P_1^{jT} \widetilde{X}_j}{P_3^{jT} \widetilde{X}_j} \right)^2 + \left(v^j - \frac{P_2^{jT} \widetilde{X}_j}{P_3^{jT} \widetilde{X}_j} \right)^2$$

here \widetilde{X} is the homogeneous representation of X . P_i^T is each row of camera projection matrix, P which is computed by $P = KR[I_{3 \times 3} - C]$. A compact representation of the rotation matrix using quaternion is a better choice to enforce orthogonality of the rotation matrix, $R = R(q)$ where q is four dimensional quaternion, i.e.,

$$\min_{C, q} \sum_{i=1, J} \left(u^j - \frac{P_1^{jT} \widetilde{X}_j}{P_3^{jT} \widetilde{X}_j} \right)^2 + \left(v^j - \frac{P_2^{jT} \widetilde{X}_j}{P_3^{jT} \widetilde{X}_j} \right)^2$$

This minimization is highly nonlinear because of the divisions and quaternion parameterization. The initial guess of the solution, (C_0, R_0) , estimated via the linear PnP is needed to minimize the cost function. This minimization can be solved using a nonlinear optimization function such as `scipy.optimize.leastsq` or `scipy.optimize.least_squares` in Scipy library.

3.7. Bundle Adjustment

Once you have computed all the camera poses and 3D points, we need to refine the poses and 3D points together, initialized by previous reconstruction by minimizing reprojection error.

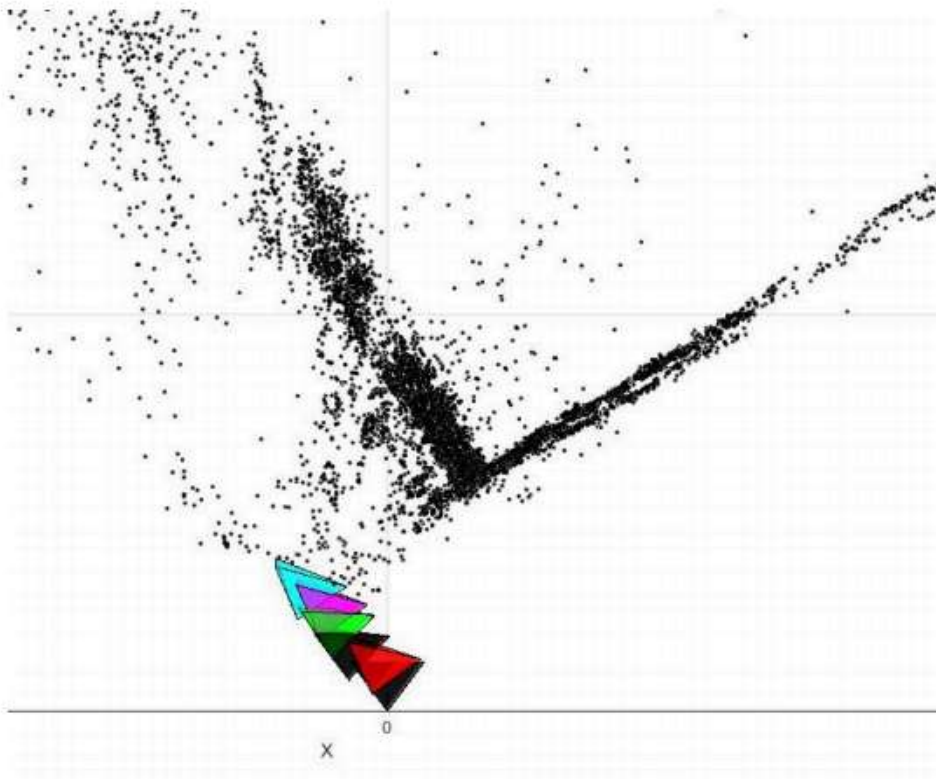


Figure 7: The final reconstructed scene after Sparse Bundle Adjustment (SBA).

3.7.1 Visibility Matrix

Find the relationship between a camera and point, construct a $I \times J$ binary matrix, V where V_{ij} is one if the j^{th} point is visible from the i^{th} camera and zero otherwise (implement the function `BuildVisibilityMatrix.py`)

3.7.2 Bundle Adjustment

Given initialized camera poses and 3D points, refine them by minimizing reprojection error (implement the function `BundleAdjustment.py`). The bundle adjustment refines camera poses and 3D points simultaneously by minimizing the following reprojection error over $C_{i=1}^I, q_{i=1}^I$ and $X_{j=1}^J$.

The optimization problem can be formulated as following:

$$\min_{\{C_i, q_i\}_{i=1}^I, \{X\}_{j=1}^J} \sum_{i=1}^I \sum_{j=1}^J V_{ij} \left(\left(u^j - \frac{P_1^{jT} \tilde{X}}{P_3^{jT} \tilde{X}} \right)^2 + \left(v^j - \frac{P_2^{jT} \tilde{X}}{P_3^{jT} \tilde{X}} \right)^2 \right) \text{ where } V_{ij} \text{ is the visibility matrix.}$$

Clearly, solving such a method to compute the structure from motion is complex and slow (*can take from several minutes for only 8-10 images*). This minimization can be solved using a nonlinear optimization functions such as `scipy.optimize.leastsq` but will be extremely slow due to a number of parameters. The Sparse Bundle Adjustment toolbox such as [pySBA](#) and [large-scale BA in scipy](#) are designed to solve such optimization by exploiting sparsity of visibility matrix, V . Note that a small number of entries in V are one because a 3D point is visible from a small subset of images. Using the sparse bundle adjustment package is not trivial and would be much faster than one you write. For SBA, you are allowed to use any optimization library.

4. Putting the pipeline together

Write a program `Wrapper.py` that run the full pipeline of structure from motion based on the above algorithm.

Also, compare your result against VSfM output. You can download the off-the-shelf SfM software here: [VSfM](#).

Project Overview

```

Data: Image Matches,  $K$ 
Result:  $X, C, R$ 
for all possible pair of images do
    | // Reject outlier correspondences
    |  $[x1, x2] = \text{GetInliersRANSAC}(x1, x2);$ 
end
// For first two images
 $F = \text{EstimateFundamentalMatrix}(x1, x2);$ 
 $E = \text{EssentialMatrixFromFundamentalMatrix}(F, K);$ 
 $[Cset, Rset] = \text{ExtractCameraPose}(E);$ 
// Perform linear triangulation
for  $i = 1:4$  do
    |  $Xseti = \text{LinearTriangulation}(K, \text{zeros}(3,1), \text{eye}(3), Cseti, Rseti, x1, x2);$ 
end
// Check cheirality condition
 $[C\ R] = \text{DisambiguateCameraPose}(Cset, Rset, Xset);$ 
// Perform Non-linear triangulation
 $X = \text{NonlinearTriangulation}(K, \text{zeros}(3,1), \text{eye}(3), C, R, x1, x2, X0);$ 
 $Cset = C, Rset = R;$ 
// Register camera and add 3D points for the rest of images
for  $i=3:I$  do
    | // Register the  $i^{\text{th}}$  image using PnP.
    |  $[Cnew\ Rnew] = \text{PnPRANSAC}(X, x, K);$ 
    |  $[Cnew\ Rnew] = \text{NonlinearPnP}(X, x, K, Cnew, Rnew);$ 
    |  $Cset = Cset \cup Cnew, Rset = Rset \cup Rnew;$ 
    | // Add new 3D points.
    |  $Xnew = \text{LinearTriangulation}(K, C0, R0, Cnew, Rnew, x1, x2);$ 
    |  $Xnew = \text{NonlinearTriangulation}(K, C0, R0, Cnew, Rnew, x1, x2, X0);$ 
    |  $X = X \cup Xnew;$ 
    | // Build Visibility Matrix.
    |  $V = \text{BuildVisibilityMatrix}(\text{traj});$ 
    | // Perform Bundle Adjustment.
    |  $[Cset\ Rset\ X] = \text{BundleAdjustment}(Cset, Rset, X, K, \text{traj}, V);$ 
end

```

Figure: The overview.

5. Notes about the Data Set

Run your SfM algorithm on the images provided [here](#). The data given to you are a set of 6 images of building in-front of Levine Hall at UPenn, using a GoPro Hero 3 with fisheye lens distortion corrected. Keypoints matching (SIFT keypoints and descriptors used) data is also provided in the same folder for pairs of images. The data folder contains 5 matching files named `matching*.txt` where `*` refers to numbers from 1 to 5. For eg., `matching3.txt` contains the matching between the third image and the fourth, fifth and sixth images, i.e.,

$\mathcal{I}_3 \leftrightarrow \mathcal{I}_4, \mathcal{I}_3 \leftrightarrow \mathcal{I}_5$ and $\mathcal{I}_3 \leftrightarrow \mathcal{I}_6$. Therefore, `matching6.txt` does not exist because it is the matching by itself.

The file format of the matching file is described next. Each matching file is formatted as follows for the i th matching file:

nFeatures: (the number of feature points of the i^{th} image - each following row specifies matches across images given a feature location in the i^{th} image.)

Each Row: (the number of matches for the j^{th} feature) (Red Value) (Green Value) (Blue Value) ($u_{\text{current image}}$) ($v_{\text{current image}}$) (image id) ($u_{\text{image id image}}$) ($v_{\text{image id image}}$) (image id) ($u_{\text{image id image}}$) ($v_{\text{image id image}}$) ...

An example of `matching1.txt` is given below:

```
nFeatures: 2002
3 137 128 105 454.740000 392.370000 2 308.570000 500.320000 4 447.580000 479.36
2 137 128 105 454.740000 392.370000 4 447.580000 479.360000
```

The images are taken at 1280×960 resolution and the camera intrinsic parameters K are given in `calibration.txt` file. You will program this full pipeline guided by the functions described in following sections.

For the extra credit: Also, capture a set of images and run your SfM algorithm. DO NOT steal images from the internet. Analyze the success and the failure of your algorithm and showcase that in your report. Note: You need to capture images, calibrate them and undistort them. Feel free to use any in-built calibration tool for this. MATLAB's calibration tool in Computer Vision toolbox will be handy.

6. Submission Guidelines

If your submission does not comply with the following guidelines, you'll be given ZERO credit

6.1. File tree and naming

Your submission on ELMS/Canvas must be a `zip` file, following the naming convention `YourDirectoryID_p3.zip`. If you email ID is `abc@umd.edu` or `abc@terpmail.umd.edu`, then your `DirectoryID` is `abc`. For our example, the submission file should be named `abc_p1.zip`. The file **must have the following directory structure** because we'll be autograding

assignments. The file to run for your project should be called `Wrapper.py`. You can have any helper functions in sub-folders as you wish, be sure to index them using relative paths and if you have command line arguments for your Wrapper codes, make sure to have default values too. Please provide detailed instructions on how to run your code in `README.md` file. Please **DO NOT** include data in your submission.

```
YourDirectoryID_hw1.zip
|  README.md
|  Your Code files
|  |  — GetInliersRANSAC.py
|  |  — EstimateFundamentalMatrix.py
|  |  — EssentialMatrixFromFundamentalMatrix.py
|  |  — ExtractCameraPose.py
|  |  — LinearTriangulation.py
|  |  — DisambiguateCameraPose.py
|  |  — NonlinearTriangulation.py
|  |  — PnP_RANSAC.py
|  |  — NonlinearPnP.py
|  |  — BuildVisibilityMatrix.py
|  |  — BundleAdjustment.py
|  |  — Wrapper.py
|  |  — Any subfolders you want along with files
|  Wrapper.py
|  Data
|  |  — BundleAdjustmentOutputForAllImage
|  |  — FeatureCorrespondenceOutputForAllImageSet
|  |  — LinearTriangulationOutputForAllImageSet
|  |  — NonLinearTriangulationOutputForAllImageSet
|  |  — PnPOutputForAllImageSetShowingCameraPoses
|  |  — Imgs/
|  — Report.pdf
```

6.2. Report

There will be no Test Set for this project. For each section of the project, explain briefly what you did, and describe any interesting problems you encountered and/or solutions you implemented. You must include the following details in your writeup:

- Please make your report extremely detailed with re-projection error after each step (Linear, Non-linear triangulation, Linear, Non-linear PnP before and after BA and so on). Describe

all the steps (anything that is not obvious) and any other observations in your report.

- Your report **MUST** be typeset in LaTeX in the IEEE Tran format provided to you in the `Draft` folder and should of a conference quality paper.
- Present the Data you collected in `Data/Imgs/`.
- Present failure cases and explanation, if any.
- Do not use any function that directly implements a part of the pipeline. If you have any doubts, please contact us via Piazza.

7. Collaboration Policy

You are encouraged to discuss the ideas with your peers. However, the code should be your own, and should be the result of you exercising your own understanding of it. If you reference anyone else's code in writing your project, you must properly cite it in your code (in comments) and your writeup. For the full honor code refer to the CMSC733 Spring 2019 website.